

Aleksandr Medvedev

LLM SOLUTIONS IMPLEMENTATION INTERN

Innopolis, Russia | Open to discussing relocation to Moscow | medv.ax@outlook.com

medvax-ai.github.io | github.com/MedvAx-AI

PROFILE

Third-year Applied AI undergraduate building reproducible LLM evaluations and ML services. Created a local Qwen2.5 routing API, measured quality and latency on 500 MMLU examples, and shipped an academic ML workflow with CI, containerized serving and model lineage. Interested in reliable LLM integration for teams.

TECHNICAL SKILLS

LLM systems: Python, Hugging Face Transformers, Qwen2.5, RAG, FAISS, experiment design, accuracy and latency evaluation

Reproducible delivery: Git, GitHub Actions CI, FastAPI, Docker Compose, DVC, MLflow, source metadata and model lineage

PROJECT EXPERIENCE

BudgetRoute-LLM | Independent project

github.com/MedvAx-AI/budgetroute-llm

- Built a local LLM routing service for a 500-example MMLU benchmark by combining small/large model selection, confidence-based cascades, abstention and retrieval-first policies with Python, Hugging Face Transformers and FastAPI.
- Benchmarked Qwen2.5-0.5B/1.5B routing on 500 paired MMLU examples by using 300 training, 100 calibration and 100 held-out questions; tracked accuracy, p50/p95 latency, route share, Brier score and ECE.
- Made semantic retrieval traceable by implementing dense embeddings and FAISS exact/HNSW search; returned document/chunk IDs, similarity scores and source metadata with each hit.

Yacht Resistance MLOps Pipeline | Academic project

github.com/MedvAx-AI/pmldl-yacht-mlops

- Made the data-to-deployment workflow reproducible and verified it with 10 automated CI tests by integrating DVC, MLflow, FastAPI, Docker Compose and GitHub Actions.
- Achieved held-out RMSE 1.280 and R^2 0.993 on 70 observations from five unseen hulls by training an Extra Trees regressor with hull-grouped train/test splits.

Beyond Proximity / SemanticSplat | Research team lead

github.com/MedvAx-AI/beyond-proximity-public

- Coordinated research tasks and experimental analysis for a five-scene, 150-query evaluation, with additional ScanNet and Replica dataset pilots.

EDUCATION

Innopolis University | BSc, Applied Artificial Intelligence | 2024–present | Third year